

Technical / Product Spec

Overview

Problem Statement: The first phase of the new client support service needs to handle onboarding issues so new clients do not get stuck during setup, access, or first-use tasks.

Product Goal: Build a support service that can detect onboarding requests, classify the blocker, guide the client through the next step, log the case, and escalate anything that cannot be resolved automatically.

Recommended Shape: Use one support orchestrator with onboarding intent detection, a small playbook library, a case tracker, and a human escalation path.

Implementation Notes: Keep this phase narrow: onboarding only, copy-ready replies only, no auto-send, and no general client support until a later roadmap item.

Goals

- Recognize onboarding-related requests quickly and separate them from general support.
- Classify the onboarding issue type and attach it to the right client case.
- Capture the minimum required context before giving instructions or escalating.
- Generate step-by-step guidance from onboarding playbooks for common blockers.
- Track case state from open to resolved so the team can see where the client is stuck.
- Escalate ambiguous or blocked cases to a human with enough context to act fast.
- Measure the most common onboarding blockers and the time it takes to resolve them.

Architecture

- An existing messaging surface receives the client request and forwards it to a support orchestrator.
- An onboarding intent classifier decides whether the message belongs in the onboarding support flow.
- A case tracker stores the client, issue type, onboarding stage, status, and ownership for every onboarding issue.
- A playbook resolver matches the issue type to a predefined onboarding playbook and drafts the next step or reply.
- A human escalation path creates a handoff task when the resolver cannot confidently finish the case.

- A logging layer stores inbound messages, drafted replies, state changes, and resolution notes for audit and review.
- A lightweight metrics layer counts issue types, reopen rates, unresolved cases, and average time to resolution.

Data model

- cases: case_id, client_id, status, issue_type, onboarding_stage, confidence, owner, created_at, updated_at, last_inbound_at, last_outbound_at, due_at, escalation_reason, resolution_summary.
- messages: message_id, case_id, direction, source_channel, sender_role, body, timestamp, external_message_id.
- playbooks: playbook_id, issue_type, required_fields, steps, fallback_questions, escalation_rules, version, last_reviewed_at.
- tasks: task_id, case_id, task_type, status, assignee, due_at, notes, created_at, updated_at.
- reminders: reminder_id, case_id, due_at, status, last_checked_at, trigger_reason.

Automation

- On inbound message, classify the request and either open a new onboarding case or attach it to the existing one.
- If the issue type is known and the required context is present, generate a copy-ready resolution message from the relevant playbook.
- If key information is missing, ask one clarifying question instead of guessing.
- If the issue is blocked, low confidence, or outside the onboarding playbooks, create a human escalation task with a summary of the conversation.
- If a case stays open past the configured reminder window, create a follow-up reminder so the team can re-engage the client.
- When the case is resolved, update the status and store the resolution summary so the same blocker can be reviewed later.

Workflows

1. New onboarding issue flow: a client sends a setup-related question, the service classifies it as onboarding, and the case tracker opens or updates the case.
2. Clarification flow: the service detects missing context, asks for the minimum needed detail, and pauses until the client replies.

3. Resolution flow: the playbook resolver returns a copy-ready answer, the team reviews it, and the case is marked resolved once the client confirms success.
4. Escalation flow: if the playbook cannot solve the issue, the service creates a human handoff with the issue type, client context, and recent messages.
5. Follow-up flow: if the client goes quiet while the case is still open, a reminder is created so the team can re-engage.

Implementation plan

1. Define the onboarding issue taxonomy and decide which issue types are in v1.
2. Collect the initial onboarding playbooks, required fields, and fallback questions for each issue type.
3. Implement the onboarding classifier and confidence threshold for routing.
4. Build the case store and state transitions for open, waiting, escalated, resolved, and closed.
5. Implement the playbook resolver so it can draft a reply from the correct onboarding steps.
6. Add the escalation handoff summary and human task creation path.
7. Add reminder logic for stalled onboarding cases and a simple way to review unresolved items.
8. Add logging and metrics so the team can see the most common onboarding blockers and the service's performance.
9. Test the flow with real sample conversations and refine the playbooks before expanding the scope.

Non-goals

- Do not support general client support categories in v1.
- Do not auto-send messages on behalf of the team.
- Do not replace a CRM or full ticketing system.
- Do not handle billing, renewals, or product feature requests in this phase.
- Do not build a public self-serve help center yet.

Acceptance criteria

- An onboarding-related message opens or updates a case within the support flow.
- The service can identify the main onboarding issue types in v1 and draft a copy-ready response for each one.
- If the service lacks confidence or needs more information, it asks a clarifying question instead of guessing.

- Blocked cases create a human handoff that includes the client, issue type, recent context, and next action.
- Every case has a visible status and a final resolution summary.
- Non-onboarding requests are not routed into this service by default.

Risks

- Short or vague requests may be misclassified without a clear taxonomy and confidence threshold.
- Incomplete onboarding playbooks may leave the service unable to resolve common blockers.
- If the support surface does not store enough client state, the service may ask for avoidable clarifying questions.
- Time-based reminders can become noisy if the case status is not updated promptly.
- A weak handoff summary can slow down human escalation and make the client wait longer.

Open questions

- Which support channel is in scope for v1: WhatsApp only, or also email or portal intake?
- What exact onboarding issue types should be covered in the first release?
- Who owns escalations once the case is handed off to a human?
- Should the service support more than one language in the first phase?
- What is the default reminder window for cases that stay unresolved?

Future ideas

- Add a visible onboarding progress tracker so the team can see the client's setup stage at a glance.
- Add proactive onboarding reminders before a client gets stuck.
- Expand the same support service to billing, renewals, and general account help after onboarding is stable.
- Add multilingual reply templates for clients who prefer Hebrew or English.
- Add analytics that show which onboarding step causes the most friction.